

Predicting Fire Burn Occurrences

Lina Nguyen, Cole Bailey, Kevin Stewart

12/3/2021

Importing packages

```
library(tidyverse);library(tidyr);library(zoo);library(zoo);library(xts);library(tseries);library(astsa);library(lubridate);library(ggplot2);library(dplyr);library(fpp2);library(dplyr);library(tsibble);library(tseries);library(forecast);library(sarima);library(fpp2);library(ggplot2)
```

```
## — Attaching packages ————— tidyverse
1.3.1 —
```

```
## ✓ ggplot2 3.3.5      ✓ purrr   0.3.4
## ✓ tibble  3.1.3      ✓ dplyr   1.0.7
## ✓ tidyr   1.1.3      ✓ stringr 1.4.0
## ✓ readr   2.0.0      ✓ forcats 0.5.1
```

```
## — Conflicts —————
tidyverse_conflicts() —
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

```
##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
```

```
##
## Attaching package: 'xts'

## The following objects are masked from 'package:dplyr':
##
##   first, last
```

```
## Registered S3 method overwritten by 'quantmod':
##   method           from
##   as.zoo.data.frame zoo
```

```
##
## Attaching package: 'lubridate'
```

```

## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union

## — Attaching packages ————— fpp2
2.4 —

## ✓ forecast 8.15      ✓ expsmooth 2.3
## ✓ fma      2.4

##

##
## Attaching package: 'fpp2'

## The following object is masked from 'package:astsa':
##
##   oil

##
## Attaching package: 'tsibble'

## The following object is masked from 'package:lubridate':
##
##   interval

## The following object is masked from 'package:zoo':
##
##   index

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, union

## Loading required package: stats4

##
## Attaching package: 'sarima'

## The following object is masked from 'package:astsa':
##
##   sarima

```

#loading the data file and seperating string values

```

df_fire <-
read.csv("/Users/kevinstewart/Desktop/business_a/san_diego_fire_incident.csv"
, sep = ",",)

```

#viewing data

```
head(df_fire)
```

```
## fd_problem_nature_agg_datasd_v1 X
X.1
## 1 agency_type address_city
problem
## 2 Fire SAN DIEGO Ringing
Alarm
## 3 Fire SAN DIEGO Cardiac / Respiratory
Arrest
## 4 Fire SAN DIEGO Assault/Rape
(L4)
## 5 Fire SAN DIEGO Assist PD - Ladder
Bldg
## 6 Fire SAN DIEGO Back Pain (Non Traumatic)
(L4)
## X.2 X.3 X.4
## 1 problem_count month_response year_response
## 2 1 1 1900
## 3 1 9 2001
## 4 2 9 2006
## 5 4 9 2006
## 6 42 9 2006
```

correct colnames & and subsetting data based on year and most occurring problem

```
names(df_fire) <- df_fire %>% slice(1) %>%
unlist()
fire <- df_fire %>% slice(-1)
```

Mutate the data to month response

```
fire$problem_count <- as.numeric(fire$problem_count)
fire$month_response <- as.numeric(fire$month_response)
fire$year_response <- as.numeric(fire$year_response)
fire$problem <- as.factor(fire$problem)

#####
# verify structure of data
str(fire)


## 'data.frame': 25138 obs. of 6 variables:
## $ agency_type : chr "Fire" "Fire" "Fire" "Fire" ...
## $ address_city : chr "SAN DIEGO" "SAN DIEGO" "SAN DIEGO" "SAN DIEGO"
...
## $ problem : Factor w/ 485 levels ".Confined Space/Trench
Rescue",...: 329 111 70 77 84 88 97 100 103 119 ...
## $ problem_count : num 1 1 2 4 42 2 9 1 1 7 ...
## $ month_response: num 1 9 9 9 9 9 9 9 9 9 ...
## $ year_response : num 1900 2001 2006 2006 2006 ...
```

Verify the length and dimensions of the data

```
#####  
# Instantiate the variable  
burns_fire <- fire %>%  
  filter(problem == "Burns / Explosion (L3)")
```


Viewing frequency of data

```
#####  
ggplot(data = fire) + geom_histogram(mapping = aes(x = month_response)) +  
  scale_x_continuous(breaks = scales::pretty_breaks(n = 12)) + ggtitle('Number  
of Fire Burn Incidences per Month') + xlab('Month') + ylab('Count')  
  
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
#####  
ggplot(data = fire) + geom_bar(mapping = aes(x = year_response)) +  
  ggtitle("Number of Fire Burn Incidences per Year") + xlab("Year") +  
  ylab("Count") + scale_x_continuous(breaks = scales::pretty_breaks(n = 16))
```

Number of Fire Burn Incidences per Year




```
#####  
ggplot(data = burns_fire) + geom_bar(mapping = aes(x = month_response)) +  
scale_x_continuous(breaks = scales::pretty_breaks(n = 12)) + ggtitle('Number  
of Fire Burn Incidences per Month') + xlab('Month') + ylab('Count')
```

Number of Fire Burn Incidences per Month



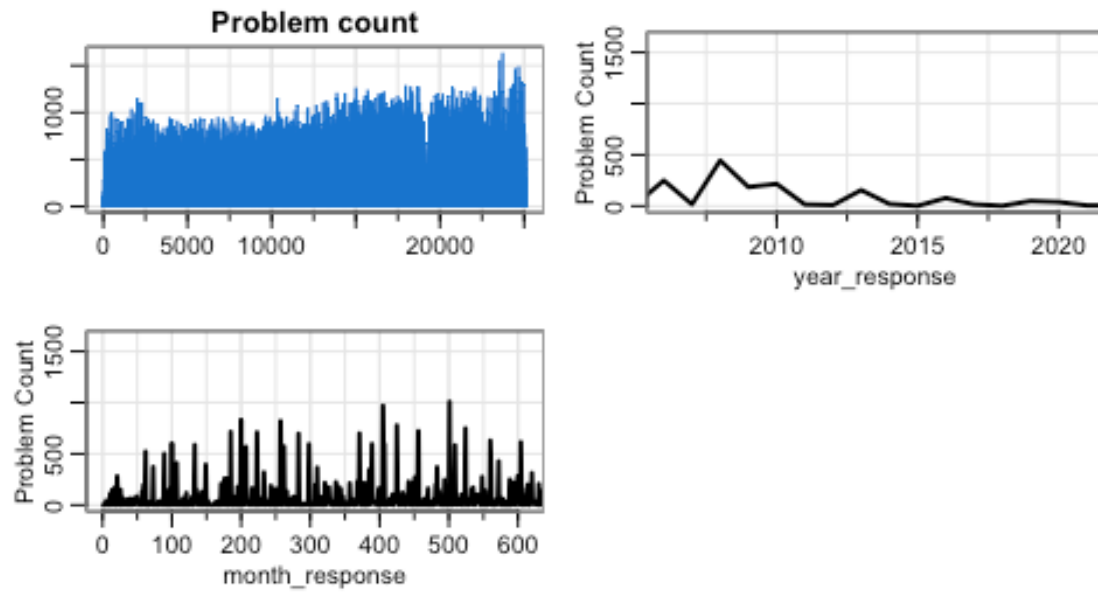
```
#####  
ggplot(data = burns_fire) + geom_bar(mapping = aes(x = year_response)) +  
scale_x_continuous(breaks = scales::pretty_breaks(n = 12)) + ggtitle('Number  
of Fire Burn Incidences per Year') + xlab('Month') + ylab('Count')
```

Number of Fire Burn Incidences per Year




Looking at time series data

```
par(mfrow=c(3:2))
tsplot(fire$problem_count, ylab = "", xlab = "", type = "l", main = "Problem
count", col = 4)
tsplot(fire[,4], xlab = "year_response", ylab = "Problem Count", type = "l",
lwd=2, xlim = c(2006, 2021))
tsplot(fire[,4], xlab = "month_response", ylab = "Problem Count", type = "l",
lwd=2, xlim = c(1,612))
```




view data by individual fire incidents

```
tsplot(burns_fire$month_response)
```

```
# logging the data graphically  
tsplot(log(burns_fire$month_response))
```



```
# the data shows that there is a there is also seasonality in the data which  
will need to be removed.
```

```
#####
```

```
#viewing first variables of the burns month response data
```


```
head(burns_fire$month_response)
```

```
## [1] 2 5 6 7 8 9
```


```
#####
```

```
# graphically displaying month response and problem count
```

```
tsplot(fire$month_response)
```



```
tsplot(fire$problem_count)
```




Declaring data as time series


```
# creating a time series variable
Y = ts(burns_fire[,4],start = c(2006,1),frequency = 12)
#####
# create a time series object to check for seasonality
ts_fire_b <- ts(burns_fire[,5],start = c(2006,1),frequency = 12)

# Checking seasonality with the seasonal graph
ggseasonplot(Y) + ggtitle("Seasonal Plot: Change in Monthly Burn Incidents ")
```

Seasonal Plot: Change in Monthly Burn Incidents



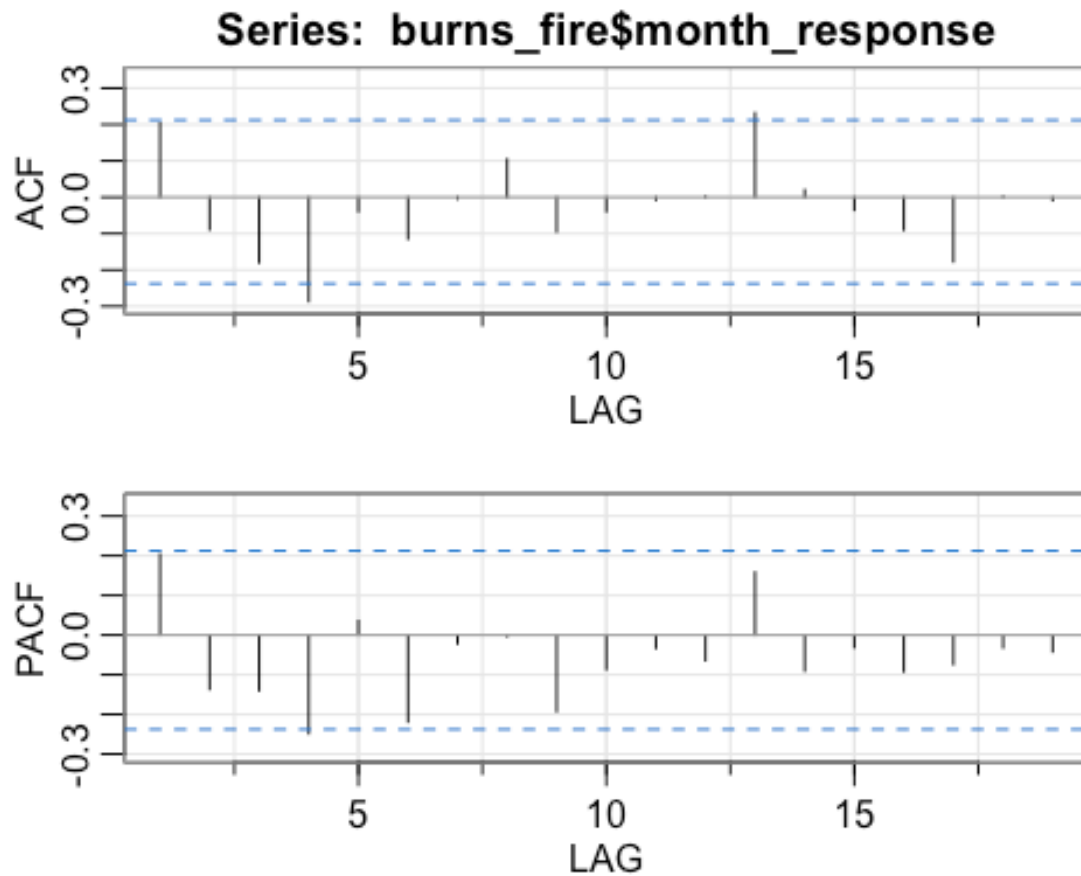
```
#####
ggsubseriesplot(ts_fire_b) +
  ggtitle("Seasonal Plot: Monthly Seasonal Graph")+ ylab("Number of Burn
Incidents")
```



#There shows that there is a spike in the winter to spring months and drastic decrease at the start of summer.

Looking at acf and pacf plot to check autocorrelation

```
head(acf2(burns_fire$month_response))
```




```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12]
## ACF   0.2 -0.09 -0.18 -0.29 -0.04 -0.12  0.00  0.1 -0.09 -0.04 -0.01  0.00
## PACF  0.2 -0.14 -0.14 -0.25  0.04 -0.22 -0.02  0.0 -0.19 -0.09 -0.03 -0.06
##      [,13] [,14] [,15] [,16] [,17] [,18] [,19]
## ACF   0.23  0.02 -0.04 -0.09 -0.18  0.00 -0.01
## PACF  0.16 -0.09 -0.03 -0.09 -0.07 -0.03 -0.04
```

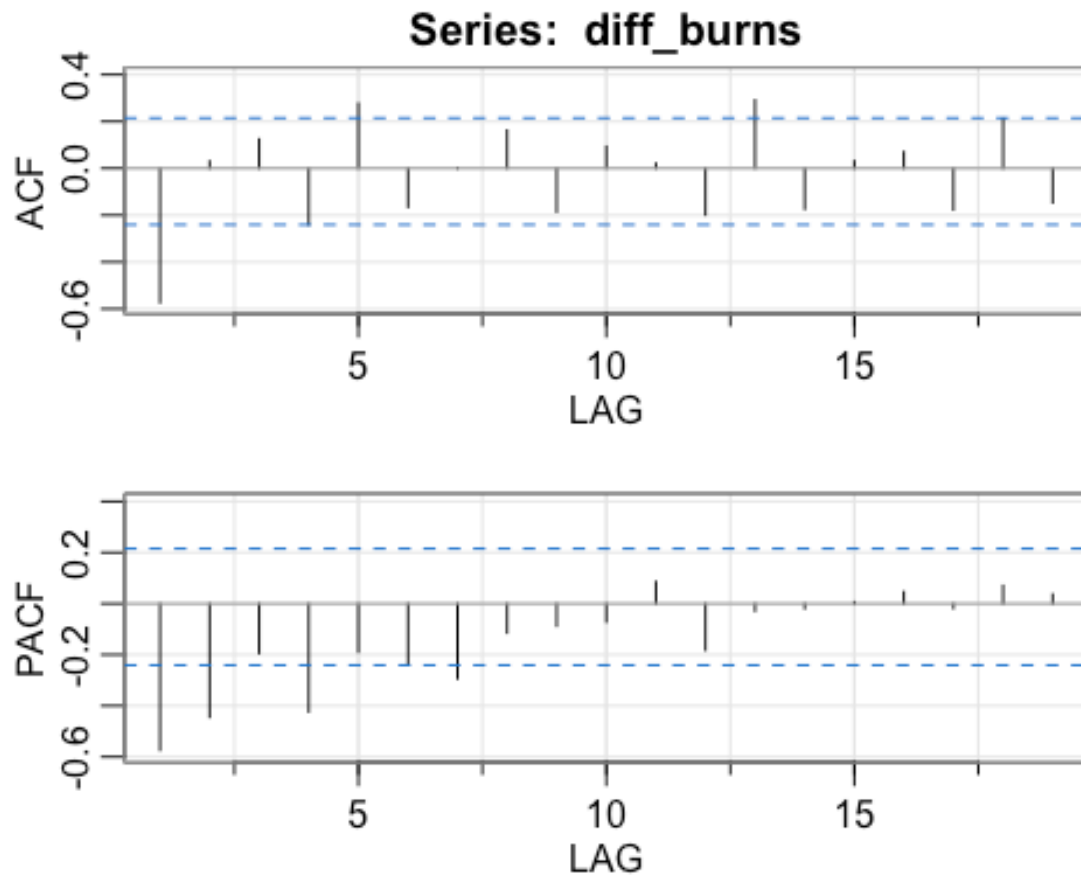
#this show a significant amount of correlation in

#Differencing the data to obtain stationarity to remove trend and seasonality

```
diff_burns <- diff(diff(burns_fire$month_response))
tsplot(diff_burns)
```



```
# acf and pacf plots  
acf2(diff_burns)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11]
##      [,12]
## ACF  -0.57  0.03  0.12 -0.24  0.28 -0.17  0.00  0.16 -0.19  0.09  0.02 -
0.20
## PACF -0.57 -0.44 -0.19 -0.42 -0.19 -0.24 -0.29 -0.11 -0.09 -0.07  0.09 -
0.18
##      [,13] [,14] [,15] [,16] [,17] [,18] [,19]
## ACF   0.29 -0.17  0.03  0.07 -0.18  0.21 -0.15
## PACF -0.03 -0.02  0.01  0.05 -0.02  0.07  0.04
```

Trying different AR models to find the best model

```
# Looking at ARIMA (1,1,1) model
ar1 <- arima(ts_fire_b, order = c(1,1,1))
ar1

##
## Call:
## arima(x = ts_fire_b, order = c(1, 1, 1))
##
## Coefficients:
##          ar1          ma1
##      0.2231  -1.0000
```

```
## s.e.  0.1128  0.0378
##
## sigma^2 estimated as 9.518:  log likelihood = -200.51,  aic = 407.03

# AIC of 262.73
#####
# ARIMA (2,0,0)
ar2 <- arima(ts_fire_b, order = c(2,0,0))
ar2

##
## Call:
## arima(x = ts_fire_b, order = c(2, 0, 0))
##
## Coefficients:
##          ar1          ar2  intercept
##          0.2350 -0.1365      6.5549
## s.e.  0.1123  0.1117      0.3792
##
## sigma^2 estimated as 9.218:  log likelihood = -199.87,  aic = 407.74

#####
# In viewing the data it appears that the data is best with differencing
ar3 <- arima(ts_fire_b, order = c(3,0,0))
ar3

##
## Call:
## arima(x = ts_fire_b, order = c(3, 0, 0))
##
## Coefficients:
##          ar1          ar2          ar3  intercept
##          0.2153 -0.1030 -0.1446      6.5708
## s.e.  0.1121  0.1134  0.1125      0.3293
##
## sigma^2 estimated as 9.022:  log likelihood = -199.05,  aic = 408.11

#####
ar4 <- arima(ts_fire_b, order = c(4,0,0))
ar4

##
## Call:
## arima(x = ts_fire_b, order = c(4, 0, 0))
##
## Coefficients:
##          ar1          ar2          ar3          ar4  intercept
##          0.1752 -0.1322 -0.0793 -0.2588      6.5735
## s.e.  0.1097  0.1110  0.1123  0.1125      0.2558
##
## sigma^2 estimated as 8.432:  log likelihood = -196.51,  aic = 405.03
```

```
#####
ar5 <- arima(ts_fire_b, order = c(5,0,0))
ar5

##
## Call:
## arima(x = ts_fire_b, order = c(5, 0, 0))
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5  intercept
##          0.1876 -0.1275 -0.0715 -0.2706  0.0498      6.5720
## s.e.    0.1132   0.1114   0.1136   0.1156   0.1162      0.2681
##
## sigma^2 estimated as 8.41:  log likelihood = -196.42,  aic = 406.85

print(ar1);ar2;ar3;ar4;ar5

##
## Call:
## arima(x = ts_fire_b, order = c(1, 1, 1))
##
## Coefficients:
##          ar1          ma1
##          0.2231 -1.0000
## s.e.    0.1128   0.0378
##
## sigma^2 estimated as 9.518:  log likelihood = -200.51,  aic = 407.03

##
## Call:
## arima(x = ts_fire_b, order = c(2, 0, 0))
##
## Coefficients:
##          ar1          ar2  intercept
##          0.2350 -0.1365      6.5549
## s.e.    0.1123   0.1117      0.3792
##
## sigma^2 estimated as 9.218:  log likelihood = -199.87,  aic = 407.74

##
## Call:
## arima(x = ts_fire_b, order = c(3, 0, 0))
##
## Coefficients:
##          ar1          ar2          ar3  intercept
##          0.2153 -0.1030 -0.1446      6.5708
## s.e.    0.1121   0.1134   0.1125      0.3293
##
## sigma^2 estimated as 9.022:  log likelihood = -199.05,  aic = 408.11
```

```
##
## Call:
## arima(x = ts_fire_b, order = c(4, 0, 0))
##
## Coefficients:
##          ar1          ar2          ar3          ar4  intercept
##          0.1752 -0.1322 -0.0793 -0.2588      6.5735
## s.e.    0.1097  0.1110  0.1123  0.1125      0.2558
##
## sigma^2 estimated as 8.432:  log likelihood = -196.51,  aic = 405.03

##
## Call:
## arima(x = ts_fire_b, order = c(5, 0, 0))
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5  intercept
##          0.1876 -0.1275 -0.0715 -0.2706  0.0498      6.5720
## s.e.    0.1132  0.1114  0.1136  0.1156  0.1162      0.2681
##
## sigma^2 estimated as 8.41:  log likelihood = -196.42,  aic = 406.85

#####
# It is determined that an ARIMA model of (4,0,0) is the best model for
forecasting
```

Find the best arima model using a fit arima

Determining best model

```
fit_ar <- auto.arima(diff_burns, stepwise = FALSE, approximation = FALSE,
trace = TRUE)
```

```
##
## ARIMA(0,0,0) with zero mean      : 506.6855
## ARIMA(0,0,0) with non-zero mean : 508.793
## ARIMA(0,0,1) with zero mean      : Inf
## ARIMA(0,0,1) with non-zero mean : Inf
## ARIMA(0,0,2) with zero mean      : Inf
## ARIMA(0,0,2) with non-zero mean : Inf
## ARIMA(0,0,3) with zero mean      : Inf
## ARIMA(0,0,3) with non-zero mean : Inf
## ARIMA(0,0,4) with zero mean      : Inf
## ARIMA(0,0,4) with non-zero mean : Inf
## ARIMA(0,0,5) with zero mean      : Inf
## ARIMA(0,0,5) with non-zero mean : Inf
## ARIMA(1,0,0) with zero mean      : 478.2234
## ARIMA(1,0,0) with non-zero mean : 480.3899
## ARIMA(1,0,1) with zero mean      : Inf
## ARIMA(1,0,1) with non-zero mean : Inf
## ARIMA(1,0,2) with zero mean      : Inf
```

```
## ARIMA(1,0,2) with non-zero mean : Inf
## ARIMA(1,0,3) with zero mean      : Inf
## ARIMA(1,0,3) with non-zero mean : Inf
## ARIMA(1,0,4) with zero mean      : Inf
## ARIMA(1,0,4) with non-zero mean : Inf
## ARIMA(2,0,0) with zero mean      : 463.2976
## ARIMA(2,0,0) with non-zero mean : 465.516
## ARIMA(2,0,1) with zero mean      : Inf
## ARIMA(2,0,1) with non-zero mean : Inf
## ARIMA(2,0,2) with zero mean      : Inf
## ARIMA(2,0,2) with non-zero mean : Inf
## ARIMA(2,0,3) with zero mean      : Inf
## ARIMA(2,0,3) with non-zero mean : Inf
## ARIMA(3,0,0) with zero mean      : 462.9781
## ARIMA(3,0,0) with non-zero mean : 465.2571
## ARIMA(3,0,1) with zero mean      : Inf
## ARIMA(3,0,1) with non-zero mean : Inf
## ARIMA(3,0,2) with zero mean      : Inf
## ARIMA(3,0,2) with non-zero mean : Inf
## ARIMA(4,0,0) with zero mean      : 448.2129
## ARIMA(4,0,0) with non-zero mean : 450.521
## ARIMA(4,0,1) with zero mean      : Inf
## ARIMA(4,0,1) with non-zero mean : Inf
## ARIMA(5,0,0) with zero mean      : 448.8007
## ARIMA(5,0,0) with non-zero mean : 451.166
##
##
## Best model: ARIMA(4,0,0) with zero mean
```

```
#####
```

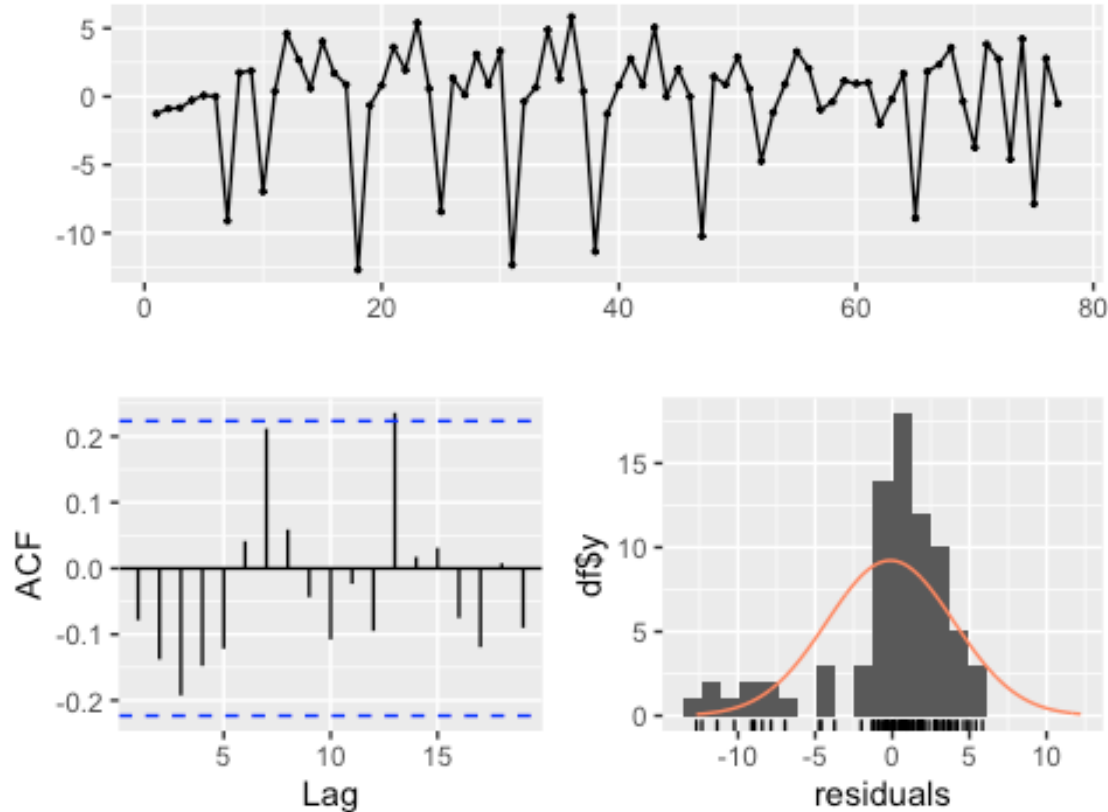
```
# Confirming that an ARIMA (4,0,0) is the best fit model
```

```
#check residuals for the best model
```

```
#check residuals using forecast and a lag
```

```
forecast::checkresiduals(fit_ar, lag=12)
```

Residuals from ARIMA(4,0,0) with zero mean



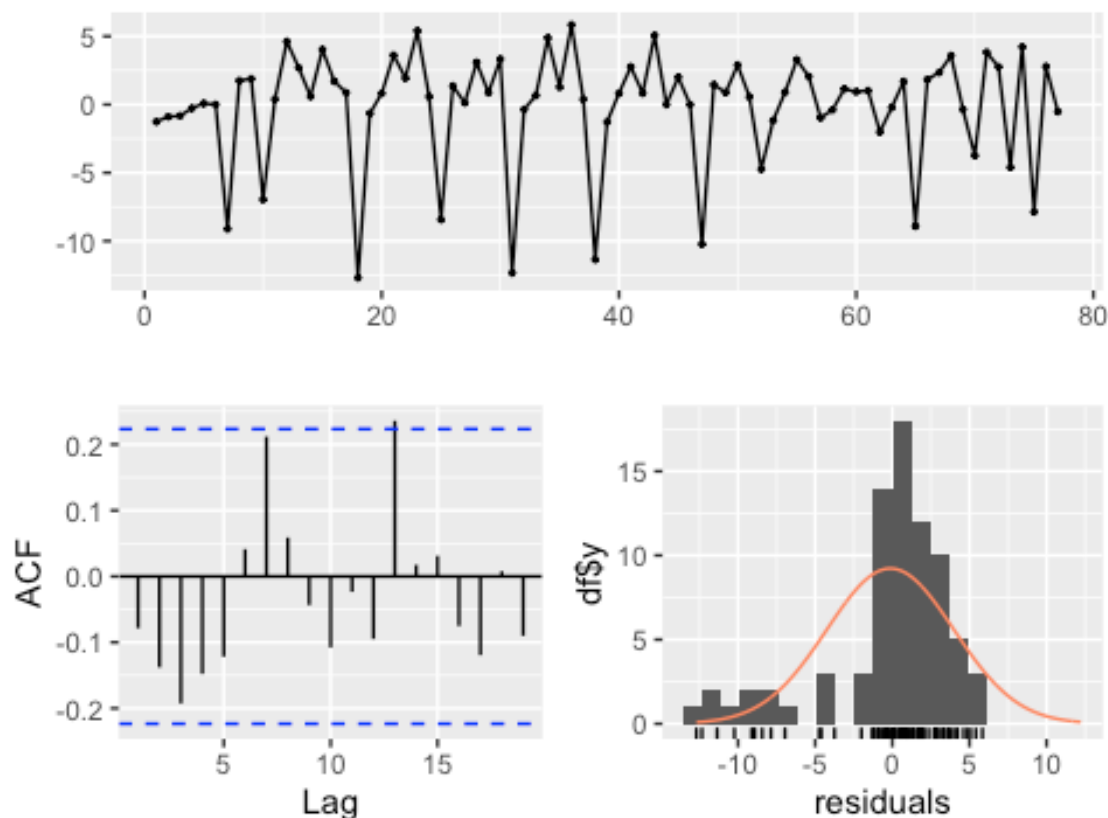
```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(4,0,0) with zero mean
## Q* = 14.637, df = 8, p-value = 0.06659
##
## Model df: 4.   Total lags used: 12

# the best model based on forecasting and checking residuals
arima(burns_fire$month_response, order = c(4,0,0))

##
## Call:
## arima(x = burns_fire$month_response, order = c(4, 0, 0))
##
## Coefficients:
##          ar1      ar2      ar3      ar4  intercept
##      0.1752 -0.1322 -0.0793 -0.2588      6.5735
## s.e.  0.1097  0.1110  0.1123  0.1125      0.2558
##
## sigma^2 estimated as 8.432:  log likelihood = -196.51,  aic = 405.03

# print summary statistics
checkresiduals(fit_ar)
```

Residuals from ARIMA(4,0,0) with zero mean



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(4,0,0) with zero mean
## Q* = 13.746, df = 6, p-value = 0.03261
##
## Model df: 4.   Total lags used: 10

print(summary(fit_ar))

## Series: diff_burns
## ARIMA(4,0,0) with zero mean
##
## Coefficients:
##          ar1      ar2      ar3      ar4
##      -0.9945  -0.8816  -0.5776  -0.4603
## s.e.   0.1009   0.1372   0.1354   0.1035
##
## sigma^2 estimated as 17.64:  log likelihood=-218.68
## AIC=447.37  AICc=448.21  BIC=459.09
##
## Training set error measures:
```

```
##               ME      RMSE      MAE MPE MAPE      MASE      ACF1
## Training set -0.1268859 4.089907 2.786639 NaN  Inf 0.3383141 -0.07872935

#Get the standard deviation
std <- sqrt(17.64)
std

## [1] 4.2
```

In viewing the acf most of all the autocorrelation is removed from the model.

Dickey fuller Test to check for stationarity

```
adf.test(diff_burns)

## Warning in adf.test(diff_burns): p-value smaller than printed p-value

##
## Augmented Dickey-Fuller Test
##
## data: diff_burns
## Dickey-Fuller = -7.0975, Lag order = 4, p-value = 0.01
## alternative hypothesis: stationary

#The stationarity is removed according to the dickey fuller test
```


We would reject the null, shows that we have stationarity.

Using arima to forecast monthly trends of fire incidents in San Diego


```
f_cast <- forecast(fit_ar, h=12)

autoplot(f_cast)
```


Forecasts from ARIMA(4,0,0) with zero mean



```
# forecasting to look at most recent data  
autoplot(f_cast, include = 48)
```



```
head(print(summary(f_cast)))
```

```
##
## Forecast method: ARIMA(4,0,0) with zero mean
##
## Model Information:
## Series: diff_burns
## ARIMA(4,0,0) with zero mean
##
## Coefficients:
##          ar1      ar2      ar3      ar4
##      -0.9945  -0.8816  -0.5776  -0.4603
## s.e.   0.1009   0.1372   0.1354   0.1035
##
## sigma^2 estimated as 17.64:  log likelihood=-218.68
## AIC=447.37  AICc=448.21  BIC=459.09
##
## Error measures:
##              ME      RMSE      MAE  MPE  MAPE      MASE      ACF1
## Training set -0.1268859 4.089907 2.786639 NaN  Inf  0.3383141 -0.07872935
##
## Forecasts:
##   Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
```

## 78	-5.06772694	-10.4508388	0.3153849	-13.300486	3.165032
## 79	7.26692522	-0.3251188	14.8589693	-4.344104	18.877955
## 80	-7.35322034	-14.9672621	0.2608214	-18.997892	4.291452
## 81	5.67460117	-2.0095082	13.3587106	-6.077230	17.426432
## 82	-1.02554519	-8.7651319	6.7140415	-12.862222	10.811131
## 83	-3.08055512	-11.1098465	4.9487363	-15.360297	9.199186
## 84	4.07472804	-4.2437698	12.3932259	-8.647317	16.796773
## 85	-3.35611585	-11.6813903	4.9691586	-16.088524	9.376293
## 86	1.99676959	-6.3644246	10.3579638	-10.790573	14.784113
## 87	0.03735123	-8.3614488	8.4361512	-12.807505	12.882207
## 88	-1.73458611	-10.1829581	6.7137859	-14.655256	11.186084
## 89	2.08358759	-6.4035933	10.5707685	-10.896435	15.063611

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## 78	-5.067727	-10.4508388	0.3153849	-13.300486	3.165032
## 79	7.266925	-0.3251188	14.8589693	-4.344104	18.877955
## 80	-7.353220	-14.9672621	0.2608214	-18.997892	4.291452
## 81	5.674601	-2.0095082	13.3587106	-6.077230	17.426432
## 82	-1.025545	-8.7651319	6.7140415	-12.862222	10.811131
## 83	-3.080555	-11.1098465	4.9487363	-15.360297	9.199186